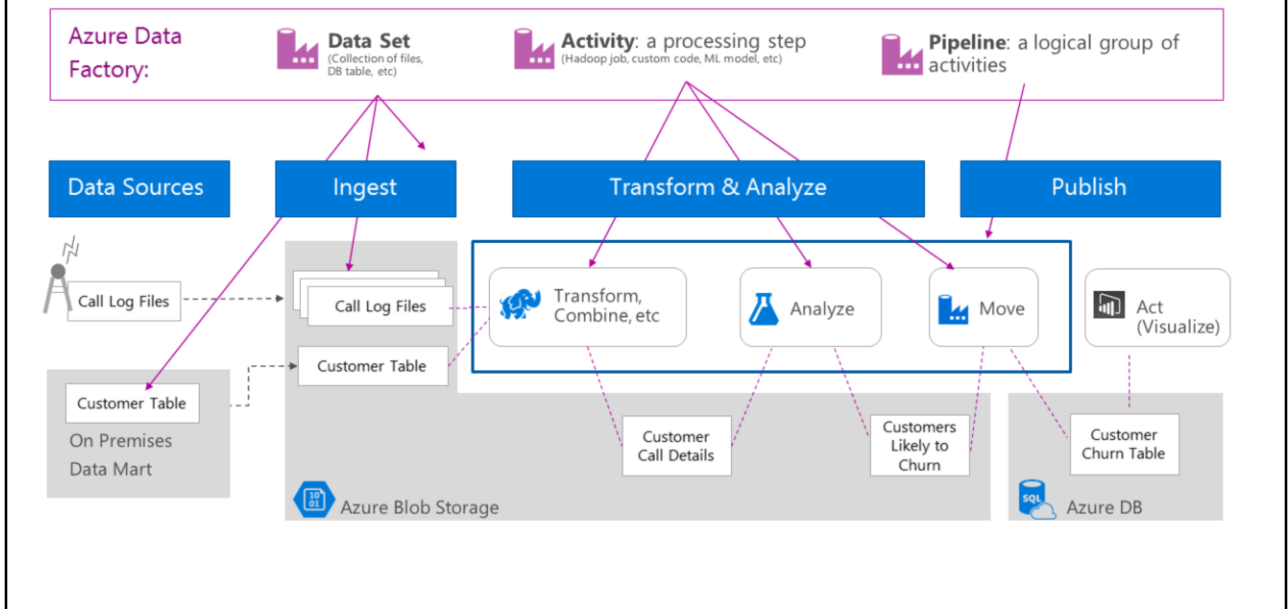




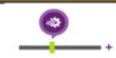








1. Main page: <http://cortanaanalytics.com>
2. Before you begin this module, you should have:
 1. Basic Math and Stats skills
 2. Business and Domain Awareness
 3. A general Computing Background
 4. A general Azure Awareness
 5. An understanding of and interpretation of the problem Statement for a Customer Scenario

Cortana Analytics Scenario - Churn



1. Full explanation of this example: <https://azure.microsoft.com/en-us/blog/getting-started-with-azure-data-factory-and-azure-machine-learning-4/>

Cortana Analytics

	Cortana
	Power BI
	Azure Stream Analytics
	Azure HDInsight
	Azure Machine Learning
	Azure SQL DB, Data Warehouse, DocumentDB
	Azure Data Lake
	Azure Event Hubs
	Azure Data Catalog
	Azure Data Factory
	Microsoft Azure

1. Platform: Microsoft Azure – <http://microsoftazure.com>
2. Storage: <https://azure.microsoft.com/en-us/documentation/services/storage/> (Blobs)
3. Azure Data Lake: <http://azure.microsoft.com/en-us/campaigns/data-lake/> (from Bing)
4. Azure Data Factory: <http://azure.microsoft.com/en-us/services/data-factory/> (Orchestration and Movement)
5. Azure Event Hubs: <http://azure.microsoft.com/en-us/services/event-hubs/> (Ingest to streaming (IoT))
6. Azure Data Catalog: <http://azure.microsoft.com/en-us/services/data-catalog/> (Social networking for data)
7. Azure DocumentDB: https://azure.microsoft.com/en-us/services/documentdb/?WT.srch=1&WT.mc_ID=SEM_JQ3fO8dU
8. Azure SQL Data Warehouse: <http://azure.microsoft.com/en-us/services/sql-data-warehouse/>
9. Azure Stream Analytics: <http://azure.microsoft.com/en-us/services/stream-analytics/> (Consume from EH, Output to Alerting (Telco Example))
10. Azure HDInsight: <http://azure.microsoft.com/en-us/services/hdinsight/> (Platform on a Platform)
11. Azure Machine Learning: <http://azure.microsoft.com/en-us/services/machine-learning/> (Prediction and Prescription - to API)
12. Revolution R Enterprise: <http://www.revolutionanalytics.com/revolution-r-enterprise> (ML)
13. Power BI: <https://powerbi.microsoft.com/>
14. Cortana: <http://blogs.windows.com/buildingapps/2014/09/23/cortana-integration-and-speech-recognition-new-code-samples/> and <https://blogs.windows.com/buildingapps/2015/08/25/using-cortana-to-interact-with-your-customers-10-by-10/>

Learning objectives

1. Understand the role of Azure Machine Learning in the Azure data analysis workflow
2. Understand feature selection options in Azure ML
3. Use Azure ML to identify high-value features



1. When you finish this Module, you will be able to:
 1. Understand the role of Azure Machine Learning in the Azure data analysis workflow
 2. Understand feature selection options in Azure ML
 3. Use Azure ML to identify high-value features

Azure Machine Learning



- Import and clean data
- Build, compare, and share models
- Deploy as an API
- Use APIs from the market

1. Home: <https://azure.microsoft.com/en-us/services/machine-learning/>
2. Gallery of Experiments: <https://gallery.cortanaanalytics.com/experiments>
3. Currently-available APIs: <https://gallery.cortanaanalytics.com/machineLearningAPIs>

Azure Machine Learning

ML Algorithms are best of breed and embrace OSS

- MS + R + Python + BYOA

ML Studio for productive development

- Faster experiments results in faster improvements
- Visual Workflows & ML Experiments at Cloud Scale

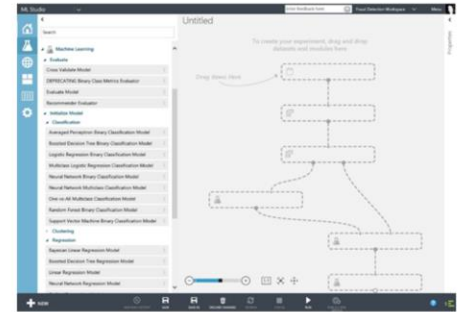
ML Operationalization to remove deployment friction

- Build entire ML Apps & Deploy as Cloud APIs

ML Applications Marketplace

- Provide ML applications like apps in an 'app store'
- Publish/consume APIs in a 2 sided market

Help organizations eliminate undifferentiated heavy lifting



Sort By: **Most Recent** Name Publisher

Customer Churn Prediction
published by: Azure Machine Learning
Customer Churn Prediction is a churn analytics service built with Azure Machine Learning. It's designed to predict the likelihood of a customer (player, subscriber, user, etc.) ending his or her relationship with a company or service.

Text Analytics
published by: Azure Machine Learning
Text Analytics API is a suite of text analytics services built with Azure Machine Learning. And bring your unstructured text (English only) and use this API to perform sentiment analysis and key phrase extraction.

Sentiment Analysis API Built with Azure Machine Learning
published by: Azure Machine Learning
This Sentiment Analysis API is an example built with Microsoft Azure Machine Learning, which assesses the sentiment polarity of short sentences, such as Facebook statuses, tweets, etc. The underlying model is built using an Azure ML, using Support Vector Machine algorithm. Based on "tweet" examples, the service outputs sentiment polarity into three levels: positive, neutral, and negative. It also provides a confidence score which could be used to further tune the polarity. The purpose of the web service is to serve as an example of how to build and publish services using Azure ML, modules, such as Train Model, Score Model, Two-class Support Vector Machine, Feature hashing, and others.

1. Home: <https://azure.microsoft.com/en-us/services/machine-learning/>
2. Gallery of Experiments: <https://gallery.cortanaanalytics.com/experiments>
3. Currently-available APIs: <https://gallery.cortanaanalytics.com/machineLearningAPIs>

Use-Cases

Sales and marketing	Finance and risk	Customer and channel	Operations and workforce
Sales forecasting	Fraud detection	User segmentation	Agent allocation
Demand forecasting	Credit risk management	Personalized offers	Warehouse efficiency
Sales lead scoring		Product recommendation	Smart buildings
Marketing mix optimization			Predictive maintenance
			Supply chain optimization

1. Collections: <http://gallery.cortanaanalytics.com/collections>
2. Real-time recommendation: <https://gallery.azureml.net/MachineLearningAPI/3574432384684cac9cc766e57729ea4c>
3. Customer churn forecasting: <https://gallery.azureml.net/MachineLearningAPI/7d86b89faf2e4cbcabb84a02179da99e>
4. Fraud detection: <https://gallery.cortanaanalytics.com/Collection/Online-Fraud-Detection-Template-1>
5. Predictive maintenance: <https://gallery.cortanaanalytics.com/Collection/Predictive-Maintenance-Template-3>
6. Perceptual Intelligence - See - <https://gallery.azureml.net/MachineLearningAPI/b0b2598aa46c4f44a08af8891e415cc7> and <https://gallery.azureml.net/MachineLearningAPI/02ce55bbc0ab4fea9422fe019995c02f> and <https://gallery.azureml.net/MachineLearningAPI/89d229231a72471ebf7280fb5bd3e18c> and <https://gallery.azureml.net/MachineLearningAPI/6948e0a54fe44e6fb70cbcc143b31298>
7. Personal Assistance - Learning, human interaction, proactive - [https://gallery.cortanaanalytics.com/browse/?categories=\["Collection"\]](https://gallery.cortanaanalytics.com/browse/?categories=[)
8. Example Video: <https://blogs.microsoft.com/business-matters/2015/07/13/dartmouth-hitchcock-ushers-in-a-new-age-of-proactive-personalized-healthcare-using-cortana-analytics-suite/>
9. Example of HowOld.net - <https://how-old.net/#>
10. Mechanics: <http://blogs.technet.com/b/machinelearning/archive/2015/05/04/fun-with-ml-stream-analytics-and-powerbi-observing-virality-in-real-time.aspx>

Types of Machine Learning Problems:



Classification: Assign a category to each item (Chinese | French | Indian | Italian | Japanese restaurant).

Regression: Predict a real value for each item (stock/currency value, temperature).

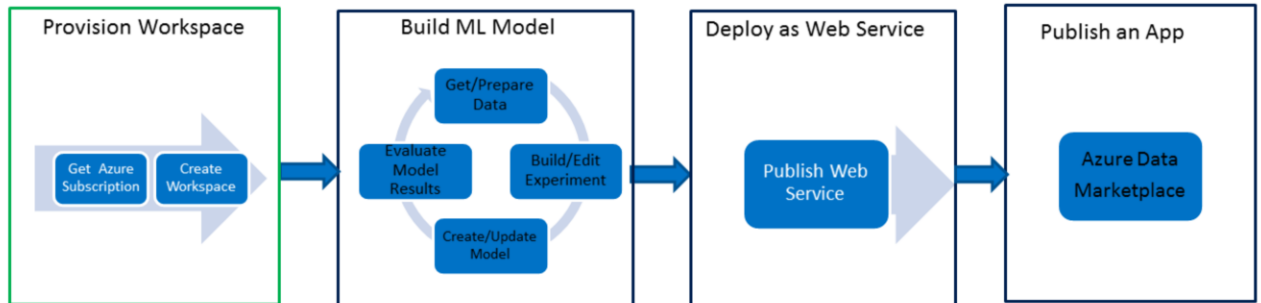
Ranking: Order items according to some criterion (web search results relevant to a user query).

Clustering: Partition items into homogeneous groups (clustering twitter posts by topic).

Dimensionality reduction: Transform an initial representation of items into a lower-dimensional representation while preserving some properties (preprocessing of digital images).

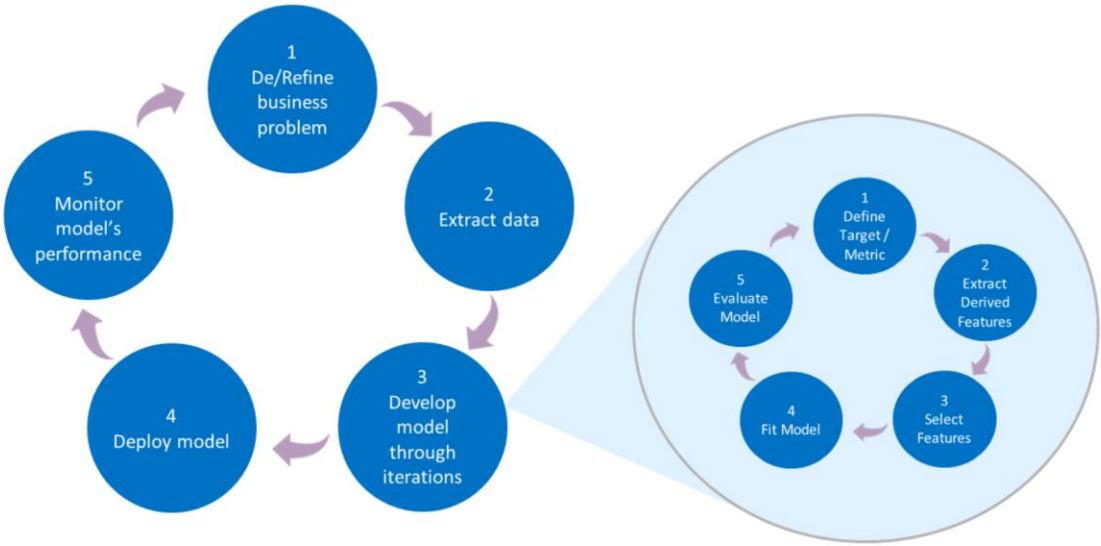
- Provided APIs: <http://gallery.cortanaanalytics.com/machineLearningAPIs>
- http://openclassroom.stanford.edu/MainFolder/courses/MachineLearning/exercises/ex8materials/ex8b_10.png
- <http://spectrum.ieee.org/img/RecommendNEWf1-1348253703748.jpg>

Azure Machine Learning Ecosystem



1. Simple overview of this process: <https://azure.microsoft.com/en-us/documentation/articles/machine-learning-create-experiment/>

Steps to Build a ML Solution




- 1. Process to learn: <https://azure.microsoft.com/en-us/documentation/learning-paths/machine-learning-self-guided-predictive-analytics-training/>

Azure Machine Learning

Data Ingress

- Read data from:
 - Web URL
 - Hive query
 - Azure SQL database
 - Azure table
 - Azure blob storage
- Load a collection of images from blob storage, for use in image classification tasks




1. <https://azure.microsoft.com/en-us/documentation/articles/machine-learning-import-data/>
2. <https://azure.microsoft.com/en-us/documentation/articles/machine-learning-data-science-preprocess-steps/>
3. <https://msdn.microsoft.com/en-us/library/azure/dn905948.aspx>
4. <https://msdn.microsoft.com/library/azure/dn905997.aspx>

Azure Machine Learning

Data Egress

- Write data out to:
 - An Azure table
 - Azure blob storage
 - Azure SQL database
 - Hive query

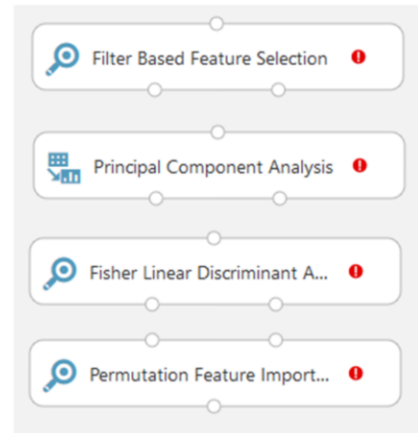


1. <https://msdn.microsoft.com/en-us/library/azure/dn905984.aspx>
2. <http://blogs.technet.com/b/machinelearning/archive/2014/11/25/azureml-web-service-parameters.aspx>

Azure Machine Learning

Feature Selection & Extraction

- Select predictive features
- Create robust new features



1. <https://azure.microsoft.com/en-us/documentation/articles/machine-learning-algorithm-choice/>
2. <https://msdn.microsoft.com/library/azure/818b356b-045c-412b-aa12-94a1d2dad90f>
3. <https://msdn.microsoft.com/library/azure/8be18eb5-ddd8-4d12-8573-7ae10d5f72fb>
4. <https://msdn.microsoft.com/library/azure/dcaab0b2-59ca-4bec-bb66-79fd23540080>
5. <https://msdn.microsoft.com/library/azure/2e010ee4-714e-44e9-933e-62d8c41818a9>

Azure Machine Learning



ML/Statistics Algorithms

- Regression
- Classification
- Clustering
- Text Analysis
- Image Analysis
- Model validation and comparison

But first: let's select features

1. Model evaluation: <https://azure.microsoft.com/en-us/documentation/articles/machine-learning-evaluate-model-performance/>

Why Use Feature Selection?



- Reduce model complexity
- Reduce overfitting
- Speed up data processing and model building

1. An Introduction to Feature Selection: <http://machinelearningmastery.com/an-introduction-to-feature-selection/>

Feature Selection in Azure ML



- Azure ML
 - Filter-based
 - Permutation Feature Importance
 - Fisher Linear Discriminant Analysis
 - Principal Components Analysis (for dimensionality reduction)
- R (caret, Boruta, etc)
- Python (scikit-learn, etc)

1. Filter-based: <https://msdn.microsoft.com/library/azure/818b356b-045c-412b-aa12-94a1d2dad90f>
2. Permutation Importance: <https://msdn.microsoft.com/library/azure/2e010ee4-714e-44e9-933e-62d8c41818a9>
3. Fisher LDA: <https://msdn.microsoft.com/library/azure/dcaab0b2-59ca-4bec-bb66-79fd23540080>
4. PCA: <https://msdn.microsoft.com/library/azure/8be18eb5-ddd8-4d12-8573-7ae10d5f72fb>
5. Comparision of LDA and PCA: http://scikit-learn.org/stable/auto_examples/decomposition/plot_pca_vs_lda.html
6. caret: <http://topepo.github.io/caret/varimp.html>
7. Boruta: <http://www.r-bloggers.com/feature-selection-all-relevant-selection-with-the-boruta-package/>
8. scikit-learn overview: <https://www.quora.com/How-do-I-perform-feature-selection>



1. Labs are real-time for this module – follow along with the instructor in the following slides

Start Azure ML Studio



- Go to: <http://studio.azureml.net/>
- Click **Get started** ➔
- If you have an Azure account, sign in; otherwise click Guest Access
- In the bottom left corner, click **+ NEW**
- Click "Blank Experiment"

Our Dataset



- Telecom data for 4,667 fictitious customers
 - Account tenure
 - Number of calls by time of day (day, evening, night)
 - Length of calls by time of day
 - Number of customer service interactions

1. <https://aka.ms/AzureMLChurnTutorial>

Read the Data into Azure ML



Data URL:


<https://telecomchurndemo.blob.core.windows.net/churnfeatures/sourcedata.csv>

1. Paste from link here: <https://telecomchurndemo.blob.core.windows.net/churnfeatures/sourcedata.csv>

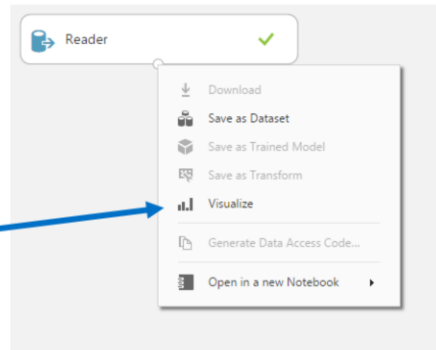
Examine the Imported Data



Hover over the output port

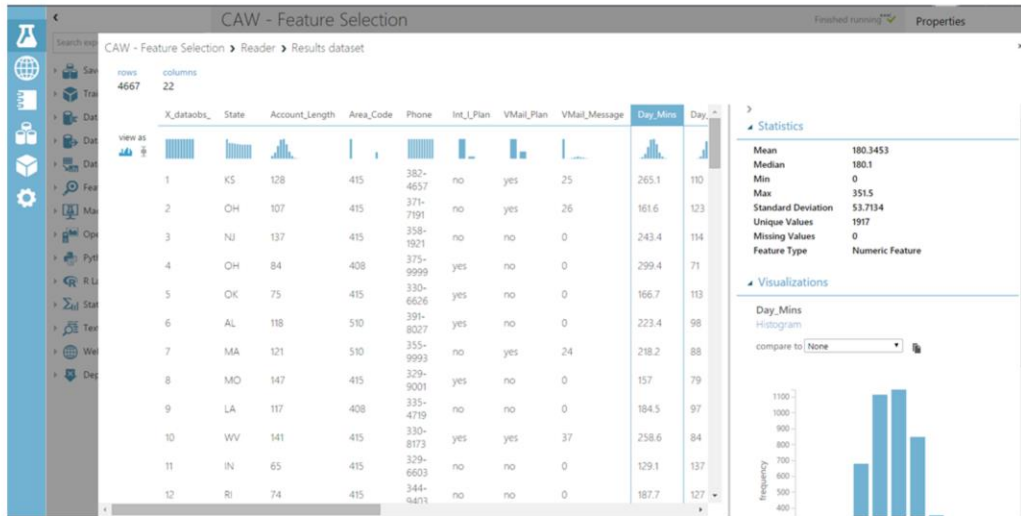


Click to see options



Click "Visualize"

Examine the Imported Data

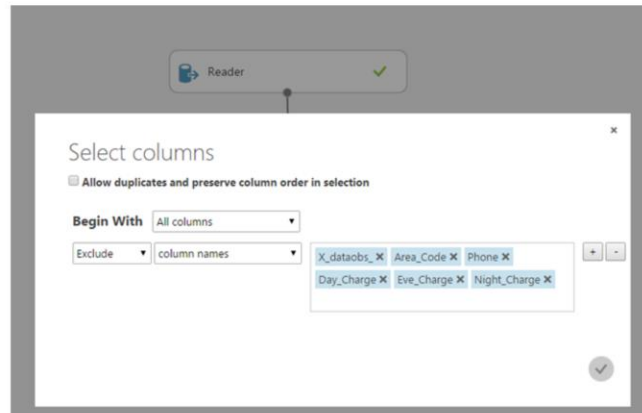
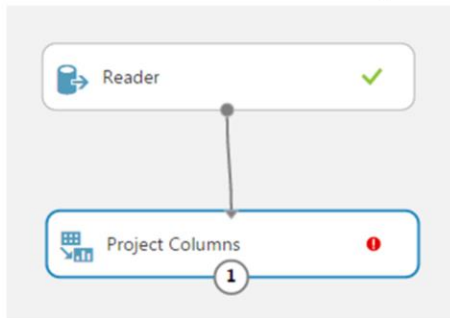


Drop Columns



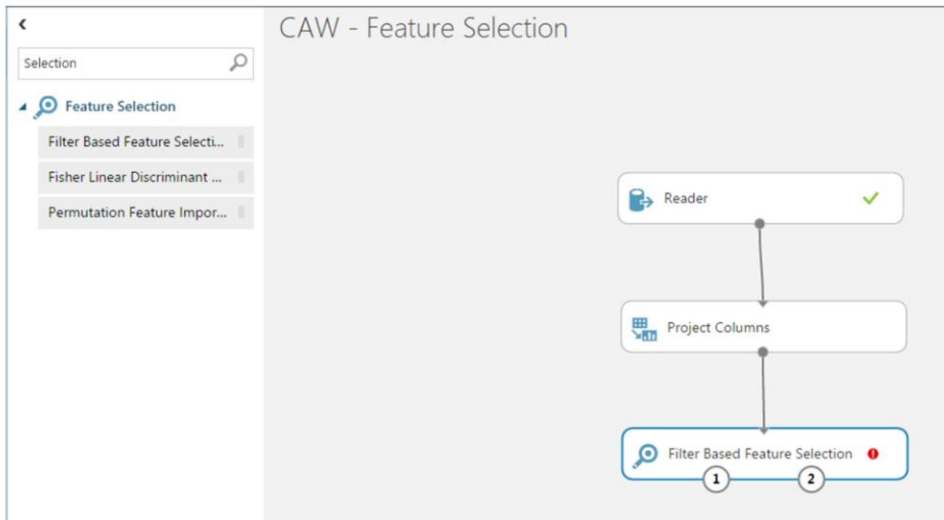
Use the Project Columns module to drop

- Non-predictors (keys, IDs, row numbers, etc)
- Highly-correlated variables (depending on your feature selection method)



- Other Data Transformation Options: <https://msdn.microsoft.com/en-us/library/azure/dn905863>
- R and Python scripts can also be used for transformations, missing-data management, etc.
 - R: <https://msdn.microsoft.com/en-us/library/azure/dn905920.aspx>
 - Python: <https://msdn.microsoft.com/en-us/library/azure/dn927167.aspx>

Select Features



1. Filter-based: <https://msdn.microsoft.com/library/azure/818b356b-045c-412b-aa12-94a1d2dad90f>
2. Fisher LDA: <https://msdn.microsoft.com/library/azure/dcaab0b2-59ca-4bec-bb66-79fd23540080>
3. Permutation Importance: <https://msdn.microsoft.com/library/azure/2e010ee4-714e-44e9-933e-62d8c41818a9>
4. Principal Components Analysis (not shown): <https://msdn.microsoft.com/library/azure/8be18eb5-ddd8-4d12-8573-7ae10d5f72fb>

Select Features



In draft

Draft saved at 1:54:40 PM

Reader

Project Columns

Filter Based Feature Selection

1 2

Properties

Filter Based Feature Selection

Feature scoring method
Pearson Correlation

☒ Operate on feature co...

Target column
Selected columns:
Column names: Churn_

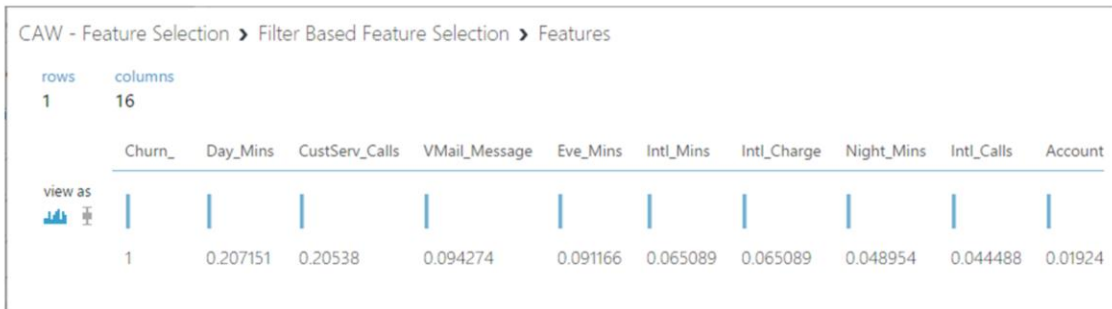
Launch column selector

Number of desired features
5

View Filter Results



- Output Port #2
- Our filter was set to use Pearson correlation – these are the raw correlations



1. We'll continue this in the next Module – make sure you save your experiment.

View Filtered Dataset



- Output Port #1
- The original dataset, now with only the top five most correlated variables
- Ready for model-building

CAW - Feature Selection > Filter Based Feature Selection > Filtered dataset

rows: 4667 columns: 6

view as:

Churn_	Day_Mins	CustServ_Calls	VMail_Message	Eve_Mins	Intl_Mins
false	265.1	1	25	197.4	10
false	161.6	1	26	195.5	13.7
false	243.4	0	0	121.2	12.2
false	299.4	2	0	61.9	6.6
false	166.7	3	0	148.3	10.1
false	223.4	0	0	220.6	6.3
false	218.2	3	24	348.5	7.5
false	157	0	0	103.1	7.1
false	184.5	1	0	351.6	8.7
false	258.6	0	37	222	11.2
true	129.1	4	0	228.5	12.7
false	187.7	0	0	163.4	9.1

Next: build a churn model, or save the output data for later use



1. Understand the role of Azure Machine Learning in the Azure data analysis workflow
2. Understand feature selection options in Azure ML
3. Use Azure ML to identify high-value features

© 2015 Microsoft Corporation. All rights reserved.

1. Use this for Q/A time