

1. Main page: <http://cortanaanalytics.com>
2. Before you begin this Module, you should have:
 1. Basic Math and Stats skills
 2. Business and Domain Awareness
 3. Familiarity with Azure ML
 4. An understanding of interpreting a problem Statement for a Customer Scenario

Learning objectives

Using Azure ML:





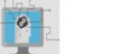






1. Split data into testing and training sets
2. Build a model for predicting customer churn
3. Compare two models
4. Share an experiment
5. Deploy a model as an API



When you have completed this Module, using Azure ML you will be able to:

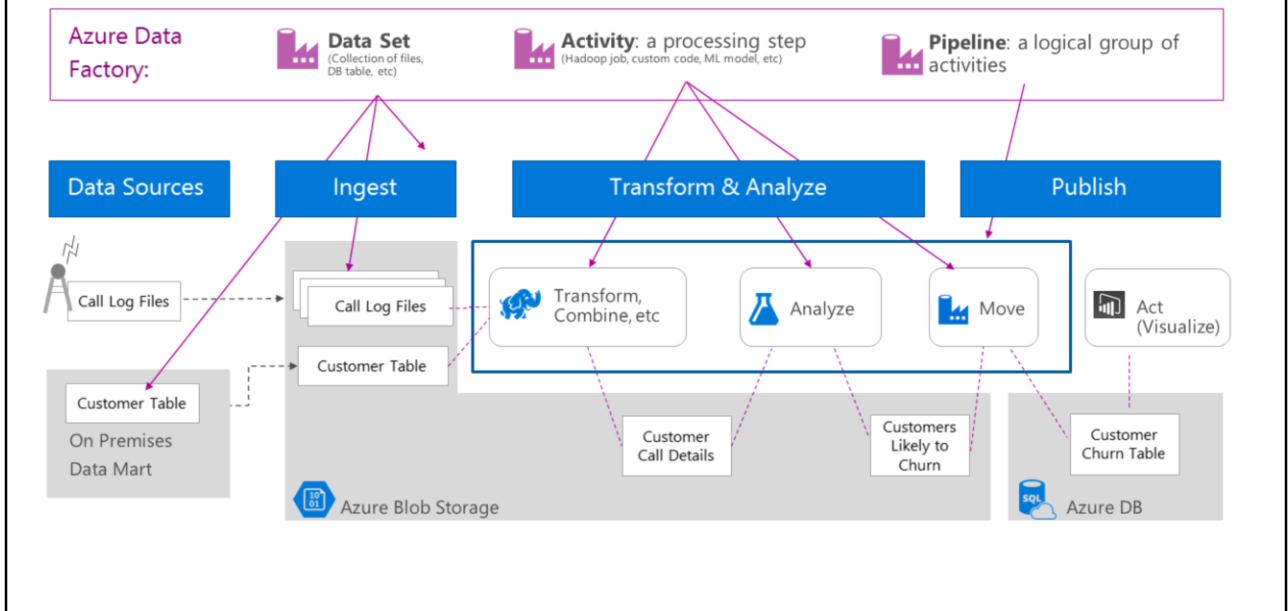
1. Split data into testing and training sets
2. Build a model for predicting customer churn
3. Compare two models
4. Share an experiment
5. Deploy a model as an API

Cortana Analytics

	Cortana
	Power BI
	Azure Stream Analytics
	Azure HDInsight
	Azure Machine Learning
	Azure SQL DB, Data Warehouse, DocumentDB
	Azure Data Lake
	Azure Event Hubs
	Azure Data Catalog
	Azure Data Factory
	Microsoft Azure

1. Platform: Microsoft Azure – <http://microsoftazure.com>
2. Storage: <https://azure.microsoft.com/en-us/documentation/services/storage/> (Blobs)
3. Azure Data Lake: <http://azure.microsoft.com/en-us/campaigns/data-lake/> (from Bing)
4. Azure Data Factory: <http://azure.microsoft.com/en-us/services/data-factory/> (Orchestration and Movement)
5. Azure Event Hubs: <http://azure.microsoft.com/en-us/services/event-hubs/> (Ingest to streaming (IoT))
6. Azure Data Catalog: <http://azure.microsoft.com/en-us/services/data-catalog/> (Social networking for data)
7. Azure DocumentDB: https://azure.microsoft.com/en-us/services/documentdb/?WT.srch=1&WT.mc_ID=SEM_JQ3fO8dU
8. Azure SQL Data Warehouse: <http://azure.microsoft.com/en-us/services/sql-data-warehouse/>
9. Azure Stream Analytics: <http://azure.microsoft.com/en-us/services/stream-analytics/> (Consume from EH, Output to Alerting (Telco Example))
10. Azure HDInsight: <http://azure.microsoft.com/en-us/services/hdinsight/> (Platform on a Platform)
11. Azure Machine Learning: <http://azure.microsoft.com/en-us/services/machine-learning/> (Prediction and Prescription - to API)
12. Revolution R Enterprise: <http://www.revolutionanalytics.com/revolution-r-enterprise> (ML)
13. Power BI: <https://powerbi.microsoft.com/>
14. Cortana: <http://blogs.windows.com/buildingapps/2014/09/23/cortana-integration-and-speech-recognition-new-code-samples/> and <https://blogs.windows.com/buildingapps/2015/08/25/using-cortana-to-interact-with-your-customers-10-by-10/>

Cortana Analytics Scenario - Churn



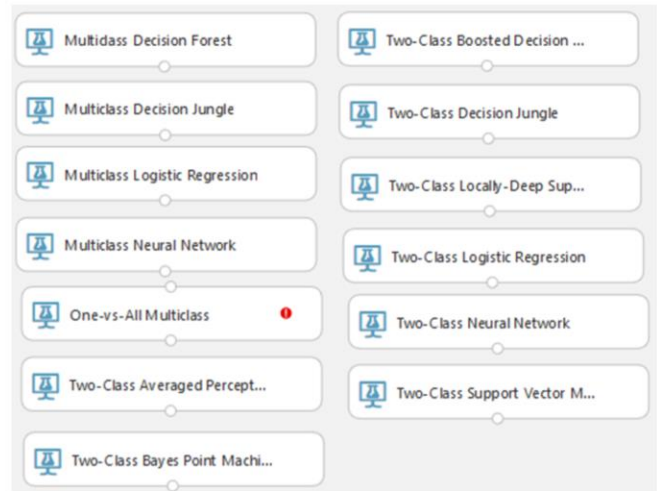
1. Full explanation of this example: <https://azure.microsoft.com/en-us/blog/getting-started-with-azure-data-factory-and-azure-machine-learning-4/>

Azure Machine Learning

Initialize Model – Classification



- Predict the class or category for data



1. Head link for this topic – note that it has a section for choosing the right algorithm:
<https://msdn.microsoft.com/en-us/library/azure/dn905812.aspx>
2. Example using flight times: <http://blogs.msdn.com/b/cdn devs/archive/2016/02/18/step-by-step-how-to-predict-the-future-with-machine-learning.aspx>

Azure Machine Learning

Initialize Model – Regression



- Algorithms that learn to predict the value of a real function for a single instance of data




1. Creating a model with a SQL Server Database:
http://blogs.msdn.com/b/business_intelligence1/archive/2014/09/14/how-to-setup-machine-learning-and-run-a-linear-regression-model-on-the-azure-cloud.aspx

Azure Machine Learning

Initialize Model – Clustering



- Create an untrained K-means clustering model




1. A good comparison of clustering versus classification:
<http://www.scientificcomputing.com/blogs/2015/09/machine-learning-techniques-help-find-patterns-big-data-sets>

Azure Machine Learning

Train



Variety of methods for training models, which can also incorporate parameter sweeping and cross-validation



1. <https://msdn.microsoft.com/library/azure/038d91b6-c2f2-42a1-9215-1f2c20ed1b40>
2. More on tuning parameters (aka hyperparameter optimization):
<http://www.johnmyleswhite.com/notebook/2012/07/21/automatic-hyperparameter-tuning-methods/>

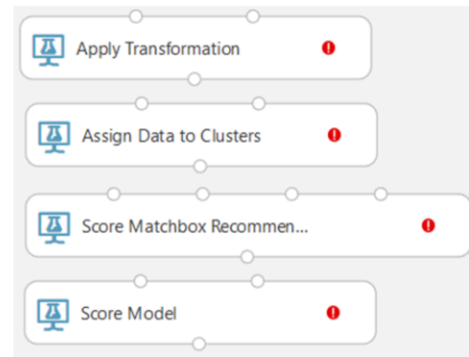
Azure Machine Learning

Score



Apply a trained model to:

- A list of recommended items
- Forecasts for time series models
- Estimates of projected demand, volume, or other numeric quantity, for regression models
- Cluster assignments
- A predicted class or outcome, for classification models
- Probability scores associated with these outputs



1. An example of training and scoring a model: <https://azure.microsoft.com/en-us/documentation/articles/machine-learning-walkthrough-4-train-and-evaluate-models/>

Azure Machine Learning

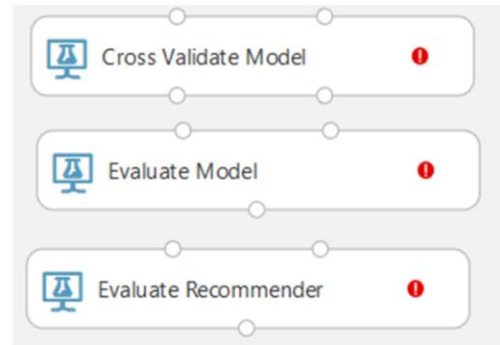
Evaluate

Metrics for Classification Models

- Accuracy
- Recall
- Precision
- F-Score
- AUC
- Average Log Loss
- Training Log Loss

Metrics for Regression Models

- Mean absolute error (MAE)
- Root mean squared error (RMSE)
- Relative absolute error (RAE)
- Relative squared error (RSE)
- Coefficient of determination



1. A discussion on using AUC effectively: <http://academic.research.microsoft.com/Paper/428291.aspx>

Azure Machine Learning

Deploy



- Access through the web service API
- Manage through the Azure management portal
- Update if the model changes

DASHBOARD CONFIGURATION

General

Published experiment

[View snapshot](#) [View latest](#)

Description

No description provided for this web service.

API key

`pw/Ce2yfwWMU23zvI9B5LCx3gQj4kloEYbWyZkZqMFXq90jKfVL8QcdHqe8Q95ERwsOpNgonu9oibzhndg =`

Default Endpoint

API HELP PAGE	TEST	APPS	LAST UPDATED
REQUEST/RESPONSE	Test	Download Excel Workbook	10/17/2015 12:26:49 PM
BATCH EXECUTION		Download Excel Workbook	10/17/2015 12:26:49 PM

1. A Video of the connection between the app and the API – we'll cover this again in the next module:
<https://channel9.msdn.com/Blogs/Windows-Azure/Introduction-to-Azure-ML-API-Service>



1. Follow the subsequent slides for the lab

Start Azure ML Studio



- Go to: <http://studio.azureml.net/>
- Click **Get started** ➔
- If you have an Azure account, sign in; otherwise click Guest Access
- In the bottom left corner, click **+ NEW**
- Click "Blank Experiment"

Our Dataset



- Telecom data for 4,667 fictitious customers
- Feature selection already complete, leaving us with:
 - Day_Mins: daytime call usage
 - Eve_Mins: evening call usage
 - Intl_Mins: international call usage
 - VMail_Message: Number of voicemails received
 - CustServ_Calls: number of customer service calls
 - Churn_: indicating whether the customer churned

1. <https://aka.ms/AzureMLChurnTutorial>

Read the Data into Azure ML



Data URL:

<https://telecomchurndemo.blob.core.windows.net/churnfeatures/churnfeatures.csv>

1. Paste from here: <https://telecomchurndemo.blob.core.windows.net/churnfeatures/churnfeatures.csv>

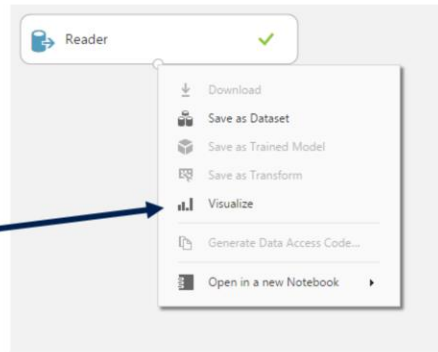
Examine the Imported Data



Hover over the output port



Right-click to see options



Click "Visualize"

Examine the Imported Data



CAW - Developing Models > Reader > Results dataset

rows
4667

columns
6

	Churn_	Day_Mins	CustServ_Calls	VMail_Message	Eve_Mins	Intl_Mins
view as						
	false	265.1	1	25	197.4	10
	false	161.6	1	26	195.5	13.7
	false	243.4	0	0	121.2	12.2
	false	299.4	2	0	61.9	6.6
	false	166.7	3	0	148.3	10.1
	false	223.4	0	0	220.6	6.3
	false	218.2	3	24	348.5	7.5
	false	157	0	0	103.1	7.1
	false	184.5	1	0	351.6	8.7
	false	258.6	0	37	222	11.2
	true	129.1	4	0	228.5	12.7
	false	187.7	0	0	163.4	9.1
	false	128.8	1	0	104.9	11.2
	false	156.6	3	0	247.6	12.3
	false	120.7	4	0	307.2	13.1
	true	332.9	4	0	317.8	5.4
	false	196.4	1	27	280.9	13.8
	false	190.7	3	0	218.2	8.1

Statistics


Mean	10.2603
Median	10.3
Min	0
Max	20
Standard Deviation	2.746
Unique Values	168
Missing Values	0
Feature Type	Numeric Feature

Visualizations

Intl_Mins

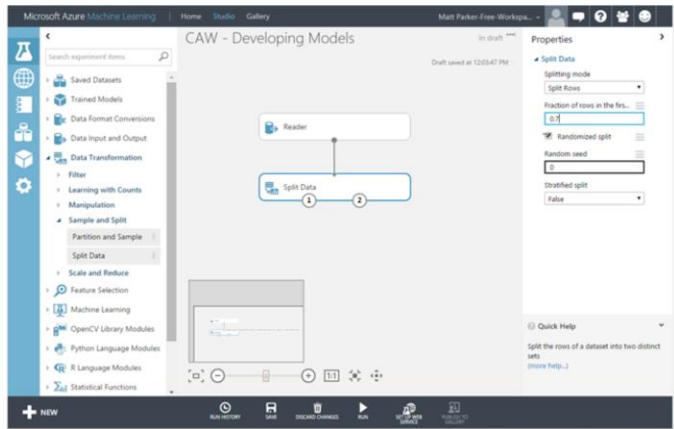
Histogram

compare to



Split the Data

- 70% training (1st port)
- 30% testing (2nd port)

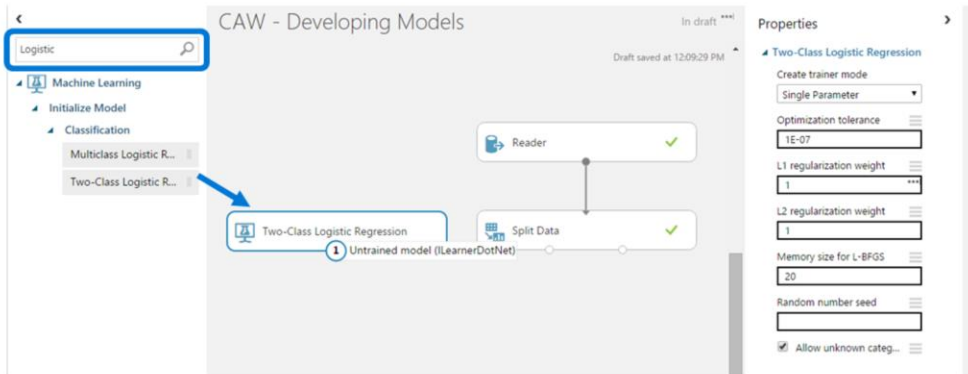


1. Splitting helps us avoid overfitting: <http://www.kdnuggets.com/2015/01/clever-methods-overfitting-avoid.html>
2. An alternative to the train-test split: Crossvalidation <https://msdn.microsoft.com/en-us/library/azure/dn905852.aspx>

Select the Model



1. Search for "Logistic" in the top left box
2. Drag "Two-Class Logistic Regression" into the experiment

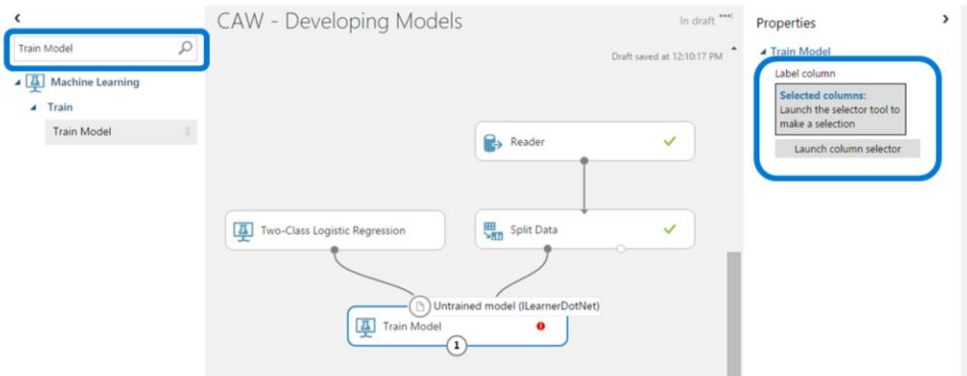


1. How to choose a model: <https://azure.microsoft.com/en-us/documentation/articles/machine-learning-algorithm-choice/>

Train the Model



1. Search for "Train Model" and drag into experiment
2. Link *untrained* model to left port, *training* data to right port

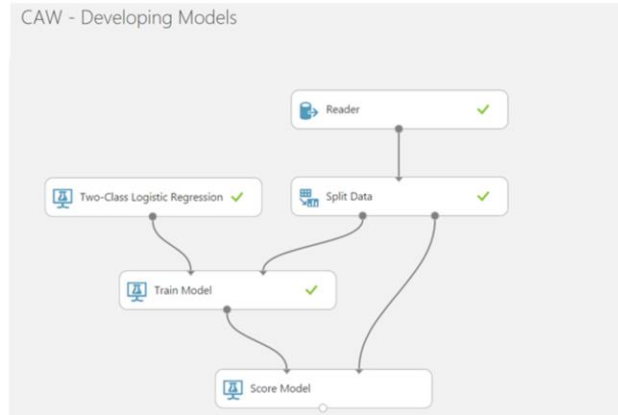


Score the Model (aka Prediction)



1. Search for "Score Model" and drag into experiment
2. Link *trained* model to left port, *testing* data to right port
3. Output: the testing dataset, *with predictions*

... but are they good predictions?



Evaluate the Model

1. Search for "Evaluate Model" and link to "Score"
2. Run experiment
3. Click the "Evaluate" out port and "Visualize"



1. Overview of evaluation statistics: https://msdn.microsoft.com/library/azure/927d65ac-3b50-4694-9903-20f6c1672089/#bkmk_classification
2. Compared with Crossvalidation: <https://azure.microsoft.com/en-us/documentation/articles/machine-learning-evaluate-model-performance/#evaluating-a-binary-classification-model>



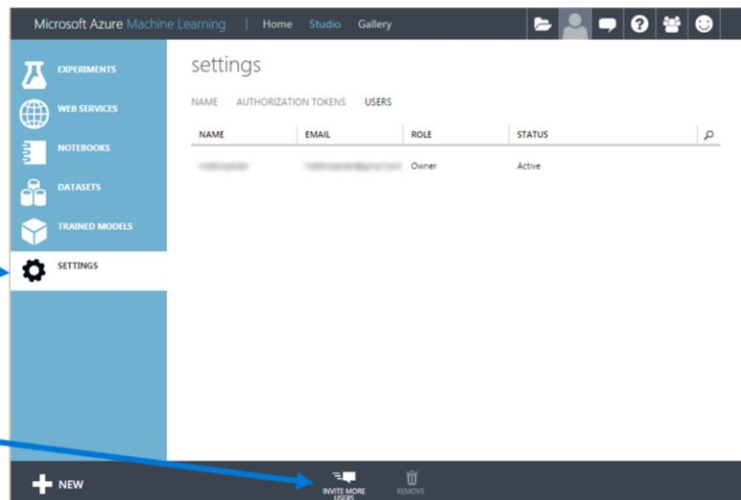
Share Your Experiment



Specifically, share your *workspace*:

1. Click the "Settings" gear on the left

2. Click "Invite More Users" at the bottom



1. More on managing your workspace: <https://azure.microsoft.com/en-us/documentation/articles/machine-learning-manage-workspace/>

Share Your Experiment



1. Enter emails of colleagues

2. Set permissions

Microsoft Azure Machine Learning | Home Studio Gallery

Invite Users to Workspace

Please specify the user's work or school account, or a Microsoft Account (e.g. Windows LiveID). The user will be notified via e-mail and will need to sign-in using that account.

Enter a work or school account, or a Microsoft Account (e.g. someone@example.com):

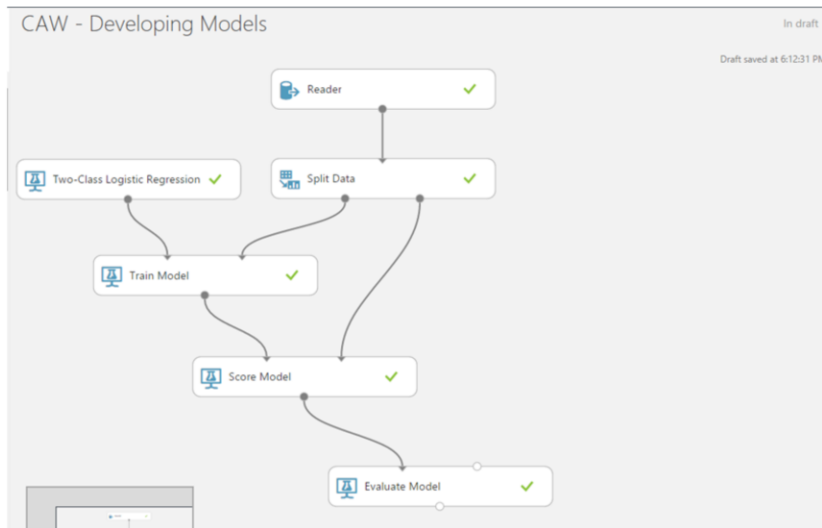
mycolleague@microsoft.com

Add as: Owner An owner can add, remove and list users with access to the workspace, in addition to what a user can do.

NEW

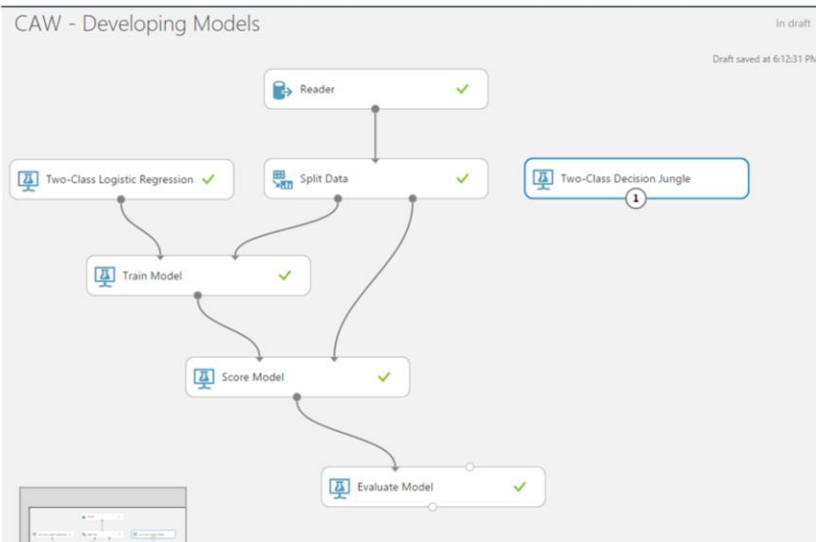
1. More on managing your workspace: <https://azure.microsoft.com/en-us/documentation/articles/machine-learning-manage-workspace/>

Add a Fancy Model



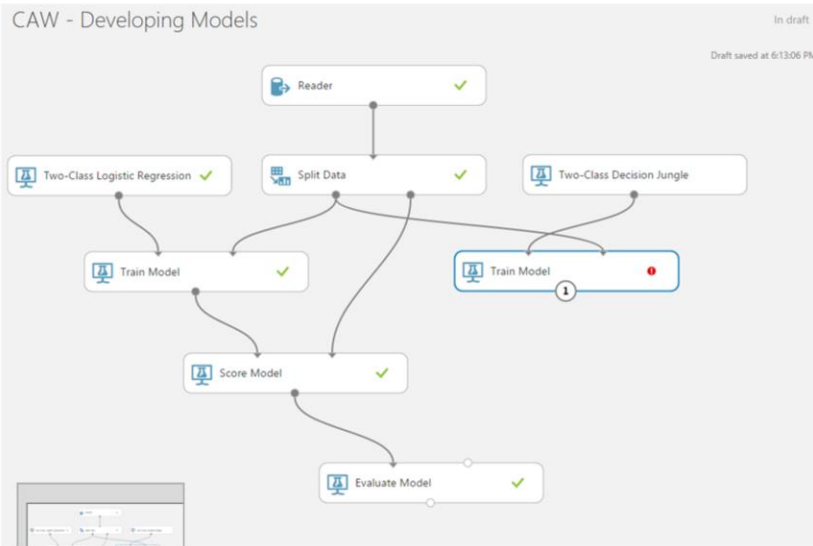
1. How to choose a model: <https://azure.microsoft.com/en-us/documentation/articles/machine-learning-algorithm-choice/>

Add a Fancy Model

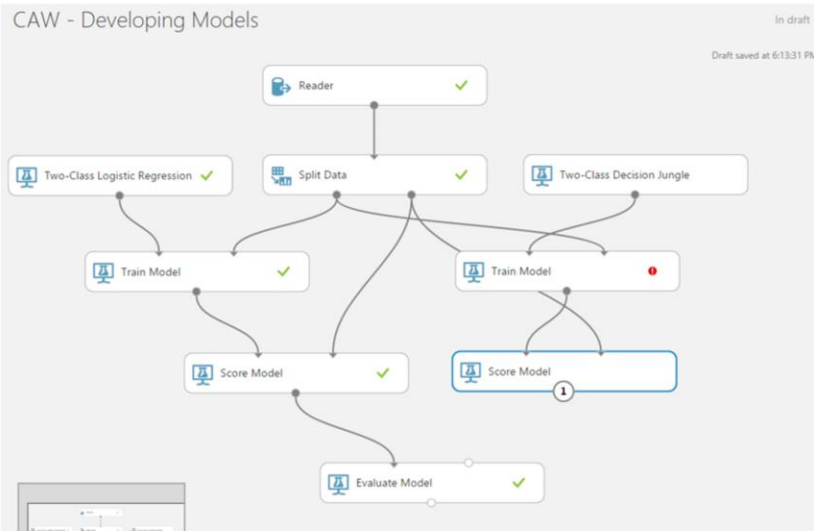


1. How to choose a model: <https://azure.microsoft.com/en-us/documentation/articles/machine-learning-algorithm-choice/>

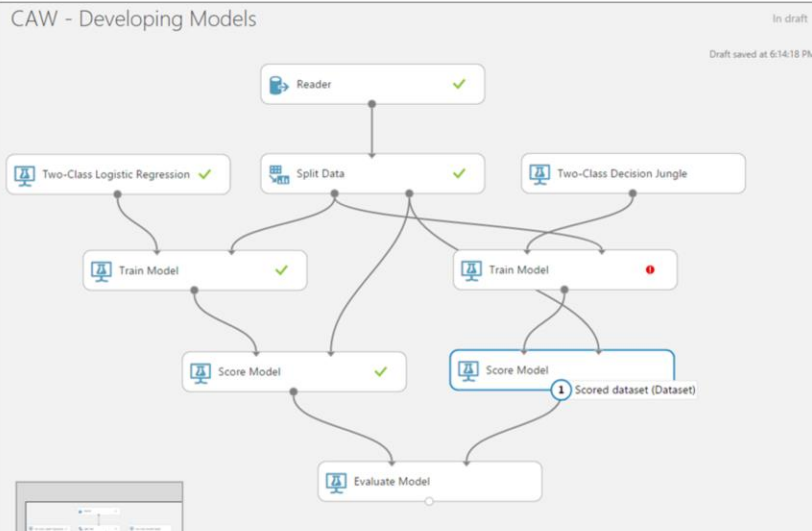
Add a Fancy Model



Add a Fancy Model




Add a Fancy Model



Compare Models

CAW - Developing Models > Evaluate Model > Evaluation results

ROC PRECISION/RECALL LIFT



True Positive	False Negative	Accuracy	Precision	Threshold	AUC
14	182	0.863	0.583	0.5	0.760
False Positive	True Negative	Recall	F1 Score		
10	1194	0.071	0.127		

1. A more complex model comparison example: <https://gallery.cortanaanalytics.com/Experiment/Compare-Binary-Classifiers-2?share=1>



Using Azure ML:

1. Split data into testing and training sets
2. Build a model for predicting customer churn
3. Compare two models
4. Share an experiment
5. Deploy a model as an API

© 2015 Microsoft Corporation. All rights reserved.

1. Use this for Q/A time